



PageRank on GCP

CONTENTS OF THIS TEMPLATE

Table of contents

01

Introduction

PageRank algorithm

03

Implementation

Solving problem manually with the theory and setups.

05

Enhancement

Compare and observations.

02

Design

Technology involved in this project

04

Test

Test with program.

06

Conclusion

Result and thoughts



01

Introduction

PageRank (PR) is an algorithm used by Google Search to rank websites in their search engine results.

It works by evaluating both the quantity and quality of links to a page.

This recursive nature of ranking forms the basis of PageRank, which considers not just the number of incoming links but also the importance of the pages providing these links, creating a reliable gauge of a page's relevance and authority on the web.

Design



02

Technology used for this project and Sections are:

GCP- Pyspark

GCP -Scala



Implementation

Three subsections involved.




STEP I: PySpark On GCP

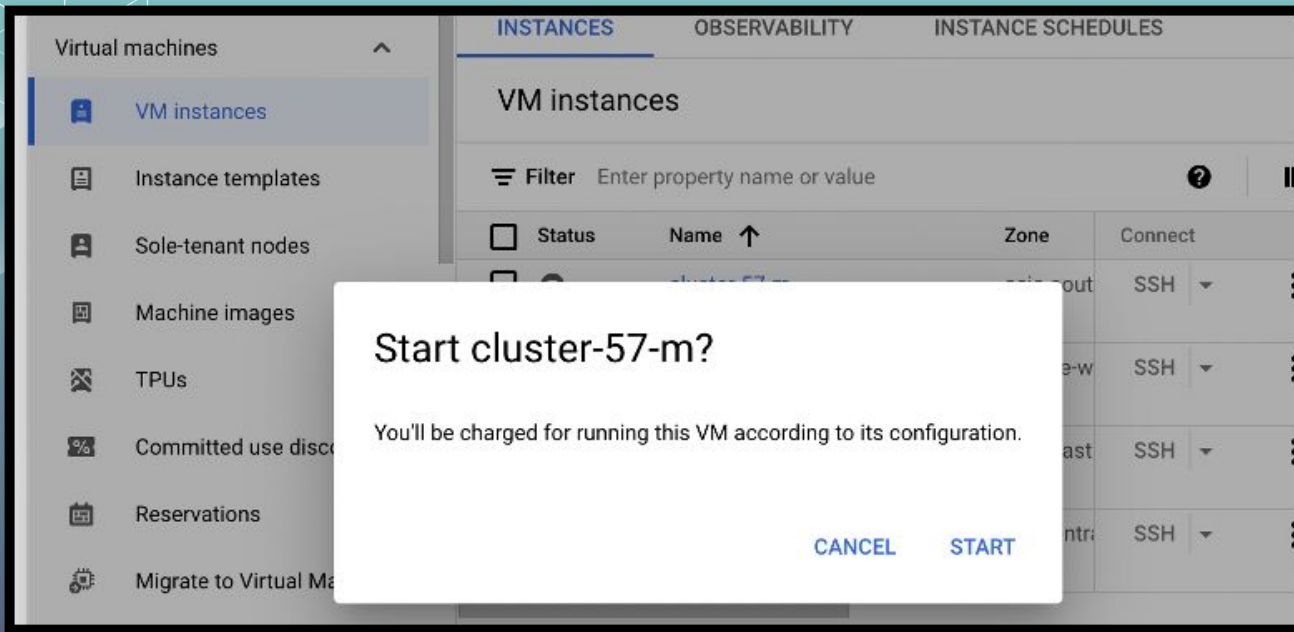
Steps of implementing PageRank with Pyspark on GCP:

- Set up environment(cluster) on GCP for Pyspark
- Create input data file for this problem
- Create Pyspark program for PageRank

Note: To run the project, need to copy the data file to HDFS file system.



I. Setup PySpark Cluster



```
Linux cluster-57-m.asia-south2-a.c.cs570-project-426508.internal 6.1.0-21-cloud-amd64 #1 SMP PREEMPT_DYNAMIC Debian 6.1.90-1 (2024-05-03) x86_64
```

```
The programs included with the Debian GNU/Linux system are free software;  
the exact distribution terms for each program are described in the  
individual files in /usr/share/doc/*/copyright.
```

```
Debian GNU/Linux comes with ABSOLUTELY NO WARRANTY, to the extent  
permitted by applicable law.
```

```
shaile32266@cluster-57-m:~$
```

2. Create Input File

```
shaile32266@cluster-57-m:~$ ls
shaile32266@cluster-57-m:~$ mkdir data
shaile32266@cluster-57-m:~$ cd data
shaile32266@cluster-57-m:~/data$ mkdir mllib
shaile32266@cluster-57-m:~/data$ cd mllib
shaile32266@cluster-57-m:~/data/mllib$ vim pagerank_data.txt
```

```
shaile32266@cluster-57-m:~/data/mllib$ vim pagerank_data.txt
shaile32266@cluster-57-m:~/data/mllib$ cat pagerank_data.txt
```

```
B C
C A
A C
A B
```

```
shaile32266@cluster-57-m:~/data/mllib$
```


2. Create PageRank PySpark Code

```
shaile32266@cluster-57-m:~/data/mllib$ cd
shaile32266@cluster-57-m:~$ mkdir PageRankProject
shaile32266@cluster-57-m:~$ cd PageRankProject
shaile32266@cluster-57-m:~/PageRankProject$ vim pagerank.py
shaile32266@cluster-57-m:~/PageRankProject$
```

For a clear visual
visit Github
repository provided
at the end of this
presentation.

```
import re
import sys
from operator import add

from pyspark.sql import SparkSession

def computeContribs(urls, rank):
    """Calculates URL contributions to the rank of other URLs."""
    num_urls = len(urls)
    for url in urls:
        yield (url, rank / num_urls)

def parseNeighbors(urls):
    """Parses a url pair string into url pair."""
    parts = re.split(':', urls)
    return parts[0], parts[1]

if __name__ == '__main__':
    if len(sys.argv) != 3:
        print("Usage: pagerank <file> <iterations>", file=sys.stderr)
        sys.exit(-1)

    print("WARN: This is a naive implementation of PageRank and is given as an example!\n" +
          "Please refer to PageRank implementation provided by graphx",
          file=sys.stderr)

    # Initialize the spark context.
    spark = SparkSession\
        .builder\
        .appName("PythonPageRank")\
        .getOrCreate()

    # Loads in input file
    lines = spark.read.text(sys.argv[1]).rdd.map(lambda r: r[0])
```

3. Create PageRank PySpark Code

```
shaile32266@cluster-57-m:~/data/mllib$ cd
shaile32266@cluster-57-m:~$ mkdir PageRankProject
shaile32266@cluster-57-m:~$ cd PageRankProject
shaile32266@cluster-57-m:~/PageRankProject$ vim pagerank.py
shaile32266@cluster-57-m:~/PageRankProject$
```

For a clear visual
visit Github
repository provided
at the end of this
presentation.

```
import re
import sys
from operator import add

from pyspark.sql import SparkSession

def computeContribs(urls, rank):
    """Calculates URL contributions to the rank of other URLs."""
    num_urls = len(urls)
    for url in urls:
        yield (url, rank / num_urls)

def parseNeighbors(urls):
    """Parses a url pair string into url pair."""
    parts = re.split(':', urls)
    return parts[0], parts[1]

if __name__ == '__main__':
    if len(sys.argv) != 3:
        print("Usage: pagerank <file> <iterations>", file=sys.stderr)
        sys.exit(-1)

    print("WARNING: This is a naive implementation of PageRank and is given as an example!\n" +
          "Please refer to PageRank implementation provided by graphx",
          file=sys.stderr)

    # Initialize the spark context.
    spark = SparkSession\
        .builder\
        .appName("PythonPageRank")\
        .getOrCreate()

    # Loads in input file
    lines = spark.read.text(sys.argv[1]).rdd.map(lambda r: r[0])
```

STEP 2: Scala On GCP

Steps of implementing PageRank with Scala on GCP:

- Set up environment(cluster) on GCP for Scala
- Create input data file for this problem
- Create Scala program for PageRank

Note: To run the project, need to copy the data file to HDFS file system.

I. Setup Scala Cluster

The screenshot shows the Google Cloud Platform console. On the left, a sidebar menu includes 'Machine images', 'TPUs', and 'Committed use discounts'. The main area displays a table of clusters. The first cluster is named 'cluster-a4bd-m' in the 'asia-southeast1-a' region, with an 'SSH' button and a three-dot menu. Below this, a terminal window titled 'SSH-in-browser' is open, showing the command prompt for the cluster. The terminal output includes the Linux distribution details and the Debian GNU/Linux warranty disclaimer.

Machine images	TPUs	Committed use discounts
		cluster-a4bd-m
		asia-southeast1-a
		SSH

SSH-in-browser

Linux cluster-a4bd-m.asia-southeast1-a.c.cs570-project-426508.internal 6.1.0-21-cloud-amd64 #1 SMP PREEMPT_DYN
MIC Debian 6.1.90-1 (2024-05-03) x86_64

The programs included with the Debian GNU/Linux system are free software;
the exact distribution terms for each program are described in the
individual files in /usr/share/doc/*/copyright.

Debian GNU/Linux comes with ABSOLUTELY NO WARRANTY, to the extent
permitted by applicable law.

shaile32266@cluster-a4bd-m:~\$

2. Install Coursier Launcher for Linux

```
shaile32266@cluster-a4bd-m:~$ curl -fL https://github.com/coursier/launchers/raw/master/cs-x86_64-pc-linux.gz |
```

```
gzip -d > cs && chmod +x cs && ./cs setup
```

```
% Total      % Received % Xferd  Average Speed   Time    Time     Time  Current
```

```
> gzip -d > cs && chmod +x cs && ./cs setup
```

% Total	% Received	% Xferd	Average Speed	Time	Time	Time	Current
			Dload	Upload	Total	Spent	Left
0	0	0	0	0	--:--:--	--:--:--	0
100	20.0M	100	20.0M	0	20.6M	--:--:--	20.6M

```
Checking if a JVM is installed
```

```
Found a JVM installed under /usr/lib/jvm/temurin-8-jdk-amd64.
```

```
Checking if ~/.local/share/coursier/bin is in PATH
```

```
Should we add ~/.local/share/coursier/bin to your PATH via ~/.profile? [Y/n] y
```

```
Checking if the standard Scala applications are installed
```

```
Installed ammonite
```

```
Installed cs
```

```
Installed coursier
```

```
Installed scala
```

```
Installed scalac
```

```
Installed scala-cli
```

```
Installed sbt
```


```
Installed sbt-n
```

```
Installed scalafmt
```

3. Set SCALA_HOME and update PATH for Scala

```
shaile32266@cluster-a4bd-m:~$ export SCALA_HOME=/usr/local/share/scala
shaile32266@cluster-a4bd-m:~$ export PATH=$PATH:$SCALA_HOME/
shaile32266@cluster-a4bd-m:~$ echo $SCALA_HOME
/usr/local/share/scala
shaile32266@cluster-a4bd-m:~$
```

```
shaile32266@cluster-a4bd-m:~$ spark-shell
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
24/06/25 12:23:18 INFO SparkEnv: Registering MapOutputTracker
24/06/25 12:23:18 INFO SparkEnv: Registering BlockManagerMaster
24/06/25 12:23:18 INFO SparkEnv: Registering BlockManagerMasterHeartbeat
24/06/25 12:23:18 INFO SparkEnv: Registering OutputCommitCoordinator
Spark context Web UI available at http://cluster-a4bd-m.asia-southeast1-a.c.cs570-project-426508.internal:33039
Spark context available as 'sc' (master = yarn, app id = application_1719315823894_0001).
Spark session available as 'spark'.
Welcome to
```



version 3.5.0

```
Using Scala version 2.12.18 (OpenJDK 64-Bit Server VM, Java 11.0.20.1)
Type in expressions to have them evaluated.
Type :help for more information.
```

```
scala>
```

Open sparkshell and copy paste this code

```
/*
 * Licensed to the Apache Software Foundation (ASF) under one or more
 * contributor license agreements.  See the NOTICE file distributed with
 * this work for additional information regarding copyright ownership.
 * The ASF licenses this file to You under the Apache License, Version 2.0
 * (the "License"); you may not use this file except in compliance with
 * the License.  You may obtain a copy of the License at
 *
 *     http://www.apache.org/licenses/LICENSE-2.0
 *
 * Unless required by applicable law or agreed to in writing, software
 * distributed under the License is distributed on an "AS IS" BASIS,
 * WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
 * See the License for the specific language governing permissions and
 * limitations under the License.
 */

// scalastyle:off println
package org.apache.spark.examples

import org.apache.spark.sql.SparkSession

/**
 * Computes the PageRank of URLs from an input file. Input file should
 * be in format of:
 * URL          neighbor URL
 * URL          neighbor URL
 * URL          neighbor URL
 * ...
 * where URL and their neighbors are separated by space(s).
 *
 * This is an example implementation for learning how to use Spark. For more conventional use,
 * please refer to org.apache.spark.graphx.lib.PageRank
 *
 * Example Usage:
 * {{{
 * bin/run-example SparkPageRank data/mllib/pagerank_data.txt 10
 * }}}
 */
```


A decorative graphic in the top-left corner consisting of a cluster of overlapping hexagons in various shades of teal and light blue, some with white outlines.

Scala Code

```
object SparkPageRank {

  def showWarning(): Unit = {
    System.err.println(
      """WARN: This is a naive implementation of PageRank and is given as an example!
        |Please use the PageRank implementation found in org.apache.spark.graphx.lib.PageRank
        |for more conventional use.
      """
    ).stripMargin()
  }

  def main(args: Array[String]): Unit = {
    if (args.length < 1) {
      System.err.println("Usage: SparkPageRank <file> <iter>")
      System.exit(1)
    }

    showWarning()

    val spark = SparkSession
      .builder
      .appName("SparkPageRank")
      .getOrCreate()

    val iters = if (args.length > 1) args(1).toInt else 10
    val lines = spark.read.textFile(args(0)).rdd
    val links = lines.map{ s =>
      val parts = s.split("\\s+")
      (parts(0), parts(1))
    }.distinct().groupByKey().cache()
    var ranks = links.mapValues(v => 1.0)
```


Scala Code

```
for (i <- 1 to iters) {  
  val contribs = links.join(ranks).values.flatMap{ case (urls, rank) =>  
    val size = urls.size  
    urls.map(url => (url, rank / size))  
  }  
  ranks = contribs.reduceByKey(_ + _).mapValues(0.15 + 0.85 * _)  
}  
  
val output = ranks.collect()  
output.foreach(tup => println(s"${tup._1} has rank: ${tup._2} ."))  
  
spark.stop()  
}  
// scalastyle:on println
```

Improved Scala Code

```
scala> val lines = sc.textFile("hdfs:///pagerank/links.txt")
lines: org.apache.spark.rdd.RDD[String] = hdfs:///pagerank/links.txt MapPartitionsRDD[1] at textFile at <console>:23

scala> val links = lines.map{ s =>
  |     val parts = s.split("\\s+")
  |     (parts(0), parts(1))
  | }.distinct().groupByKey().cache()
links: org.apache.spark.rdd.RDD[(String, Iterable[String])] = ShuffledRDD[6] at groupByKey at <console>:26

scala> var ranks = links.mapValues(v => 1.0)
ranks: org.apache.spark.rdd.RDD[(String, Double)] = MapPartitionsRDD[7] at mapValues at <console>:23

scala> for (i <- 1 to 2) {
  |     val contribs = links.join(ranks).values.flatMap{ case (urls, rank) =>
  |         val size = urls.size
  |         urls.map(url => (url, rank / size))
  |     }
  |     ranks = contribs.reduceByKey(_ + _).mapValues(0.15 + 0.85 * _)
  |     println("Iteration --> " + i)
  |     val output = ranks.collect()
  |     output.foreach(tup => println(s"${tup._1} has rank:  ${tup._2} ."))
  | }
```



Testing

Run the program with input data file to test.



Pyspark: Create HDFS file system and copy input data file from local to hdfs:

```
shaile32266@cluster-57-m:~/PageRankProject$ cd
shaile32266@cluster-57-m:~$ hdfs dfs -mkdir hdfs:///data
shaile32266@cluster-57-m:~$ hdfs dfs -mkdir hdfs:///data/mllib
shaile32266@cluster-57-m:~$ hdfs dfs -put ./data/mllib/* hdfs:///data/mllib
shaile32266@cluster-57-m:~$ hdfs dfs -ls hdfs:///data/mllib
Found 1 items
-rw-r--r--    1 shaile32266 hadoop           17 2024-06-25 12:31 hdfs:///data/mllib/pagerank_data.txt
shaile32266@cluster-57-m:~$
```

Test: Pyspark

```
shaile32266@cluster-57-m:~$ spark-submit PageRankProject/pagerank.py hdfs:///data/mllib/pagerank_data.txt 2
```

```
WARN: This is a naive implementation of PageRank and is given as an example!
```

```
Please refer to PageRank implementation provided by graphx
```

```
24/06/25 12:46:06 INFO SparkEnv: Registering MapOutputTracker
```

```
24/06/25 12:46:06 INFO SparkEnv: Registering BlockManagerMaster
```

```
24/06/25 12:46:06 INFO SparkEnv: Registering BlockManagerMasterHeartbeat
```

```
24/06/25 12:46:07 INFO SparkEnv: Registering OutputCommitCoordinator
```

```
24/06/25 12:46:07 INFO DefaultNoHARMPFailoverProxyProvider: Connecting to ResourceManager at cluster-57-m.asia-south2-a.c.cs570-project-426508.internal./10.190.0.8:8032
```

```
24/06/25 12:46:07 INFO AHSPProxy: Connecting to Application History server at cluster-57-m.asia-south2-a.c.cs570-project-426508.internal./10.190.0.8:10200
```

```
24/06/25 12:46:09 INFO Configuration: resource-types.xml not found
```

```
24/06/25 12:46:09 INFO ResourceUtils: Unable to find 'resource-types.xml'.
```

```
24/06/25 12:46:10 INFO YarnClientImpl: Submitted application application_1719313894375_0003
```

```
24/06/25 12:46:11 INFO DefaultNoHARMPFailoverProxyProvider: Connecting to ResourceManager at cluster-57-m.asia-south2-a.c.cs570-project-426508.internal./10.190.0.8:8032
```

```
24/06/25 12:46:12 INFO MetricsConfig: Loaded properties from hadoop-metrics2.properties
```

```
24/06/25 12:46:12 INFO MetricsSystemImpl: Scheduled Metric snapshot period at 10 second(s).
```

```
24/06/25 12:46:12 INFO MetricsSystemImpl: google-hadoop-file-system metrics system started
```

```
24/06/25 12:46:13 INFO GoogleCloudStorageImpl: Ignoring exception of type GoogleJsonResponseException; verified object already exists with desired state.
```

```
24/06/25 12:46:14 INFO GoogleHadoopOutputStream: hflush(): No-op due to rate limit (RateLimiter[stableRate=0.2qps]): readers will *not* yet see flushed data for gs://dataproc-temp-asia-south2-1004025664060-6rxixkt/5eac23b5-9edb-4dc7-b57f-72aebalc3a0f/spark-job-history/application_1719313894375_0003.inprogress [CONTEXT ratelimit_period="1 MINUTES" ]
```

```
-----  
Iteration --> 1
```

```
-----  
Iteration --> 1
```

```
C has rank: 1.4249999999999998.
```

```
A has rank: 1.0.
```

```
B has rank: 0.575.
```

```
-----  
Iteration --> 2
```

```
-----  
C has rank: 1.06375.
```

```
B has rank: 0.575.
```

```
A has rank: 1.3612499999999996.
```

Test: Scala

I. Create file: links.txt

```
shaile32266@cluster-a4bd-m:~$ vi links.txt
shaile32266@cluster-a4bd-m:~$ cat links.txt
B C
C A
A C
A B
shaile32266@cluster-a4bd-m:~$
```

2. Copy input file from local to HDFS

```
shaile32266@cluster-a4bd-m:~$ hdfs dfs -put links.txt .
shaile32266@cluster-a4bd-m:~$ hdfs dfs -ls
Found 2 items
drwxr-xr-x  - shaile32266 hadoop          0 2024-06-25 12:23 .sparkStaging
-rw-r--r--  1 shaile32266 hadoop       16 2024-06-25 13:00 links.txt
shaile32266@cluster-a4bd-m:~$
```

Test: Scala

3. Run-example SparkPageRank links.txt 2

```
shaile32266@cluster-a4bd-m:~$ run-example SparkPageRank links.txt 2
WARN: This is a naive implementation of PageRank and is given as an example!
Please use the PageRank implementation found in org.apache.spark.graphx.lib.PageRank
for more conventional use.
```

```
24/06/25 13:05:36 INFO SparkEnv: Registering MapOutputTracker
24/06/25 13:05:36 INFO SparkEnv: Registering BlockManagerMaster
24/06/25 13:05:36 INFO SparkEnv: Registering BlockManagerMasterHeartbeat
24/06/25 13:05:36 INFO SparkEnv: Registering OutputCommitCoordinator
```

4. Result

```
B has rank: 0.575 .
A has rank: 1.3612499999999996 .
C has rank: 1.06375 .
```


5. Test in Spark Shell with Modified Code

```
shaile32266@cluster-a4bd-m:~$ hdfs dfs -mkdir hdfs:///pagerank
shaile32266@cluster-a4bd-m:~$ hdfs dfs -put links.txt hdfs:///pagerank
shaile32266@cluster-a4bd-m:~$ hdfs dfs -ls hdfs:///pagerank
Found 1 items
-rw-r--r--  1 shaile32266 hadoop          16 2024-06-25 13:18 hdfs:///pagerank/links.txt
shaile32266@cluster-a4bd-m:~$
```


The modified code

```
scala> val lines = sc.textFile("hdfs:///pagerank/links.txt")
lines: org.apache.spark.rdd.RDD[String] = hdfs:///pagerank/links.txt MapPartitionsRDD[1] at textFile at <console>:23

scala> val links = lines.map{ s =>
  |   val parts = s.split("\\s+")
  |   (parts(0), parts(1))
  | }.distinct().groupByKey().cache()
links: org.apache.spark.rdd.RDD[(String, Iterable[String])] = ShuffledRDD[6] at groupByKey at <console>:26

scala> var ranks = links.mapValues(v => 1.0)
ranks: org.apache.spark.rdd.RDD[(String, Double)] = MapPartitionsRDD[7] at mapValues at <console>:23

scala> for (i <- 1 to 2) {
  |   val contribs = links.join(ranks).values.flatMap{ case (urls, rank) =>
  |     val size = urls.size
  |     urls.map(url => (url, rank / size))
  |   }
  |   ranks = contribs.reduceByKey(_ + _).mapValues(0.15 + 0.85 * _)
  |   println("Iteration --> " + i)
  |   val output = ranks.collect()
  |   output.foreach(tup => println(s"${tup._1} has rank:  ${tup._2} ."))
  | }
```



Enhancement



05

- **Integrate Google Cloud Storage (GCS) for data management.**
- **Optimize Spark configurations for better performance.**
- **Implement advanced data partitioning strategies.**
- **Explore real-time data processing with Apache Kafka.**
- **Create interactive dashboards for PageRank visualization.**
- **Utilize GCP's autoscaling for resource management**

Conclusion



06

- Implemented PageRank with PySpark and Scala on GCP.
- Gained practical experience in cloud-based big data processing.
- Achieved efficient handling of large datasets.
- Demonstrated scalability and high performance on GCP.
- Enhanced skills in cloud computing and data analytics.
- Ready to apply techniques to broader data projects.



GITHUB LINK

<https://github.com/Sharon20222/Cloud-Computing/tree/main/Spark/PageRank>



References

- <https://www.scala-lang.org/>
- <https://spark.apache.org/docs/latest/api/python/index.html>